

CS 520: Assignment 1

Fast Trajectory Replanning

Swapnil Shrikrishna Verlekar (sv725)
Sanket Dattakumar Dalvi (sd1482)
Saurabh Sunil Kamble (sk2675)

October 15, 2022

Part 0

We have generated 50 grid worlds of size 101×101 in java and stored them using 2-D Array.




Figure 1: Sample grid world for A* algorithm

Part 1

a)

In the example below (fig 8), the agent starts with an assumption that there are no blockers in the maze, and based on that; the agent tries to find the shortest path to the goal. Since the agent finds the goal to be two blocks away to the east, it keeps moving towards the east, observing its neighbors till it encounters a wall in its path. The agent then runs the algorithm again, considering the obstacles observed on the way to the current cell. For our implementation, the neighbors to the east, west, and north are explored when the agent starts exploring from the starting position. It then puts the neighbors in the open list and calculates the $f(n)$ for each of the neighbors, which is $g(n) + h(n)$. The heuristic value is calculated based on the Manhattan distance from the current position to the goal position. Agent then traverses to the neighbor with the least $f(n)$, which is calculated to be the East neighbor. So, in this case, according to the diagram, the costs for all directions are as below.

For West: $f(n) = g(n) + h(n)$; $f(n) = 1 + 4$; $f(n)$ from the West block is 5.
For North: $f(n) = g(n) + h(n)$; $f(n) = 1 + 4$; $f(n)$ from the South block is 5.
For East: $f(n) = g(n) + h(n)$; $f(n) = 1 + 2$; $f(n)$ from the East block is 3.
So, as $f(n)$ for the east block is the lowest among all, the first move of agent will be east rather than in other directions.

b)

For our algorithm, we add the agents current position to our closed list (contains all the nodes we have expanded) so that it will not be visited again. Then we begin by exploring neighbors from the agent's current position. We check whether all the neighbors are inside the wall (within the bound) and are not blockers. If it is not a blocker and not present in the closed list then we add them to our open list (priority queue implemented using heap). Otherwise, if neighbor is already present in the open list we will update its corresponding values $f(n)$, $g(n)$, $h(n)$; if the new $f(n)$ is lesser than the previous $f(n)$, then same process is repeated for the its neighbors. For example a 100x100 maze with have $1 \leq NodesTravelled \leq 10000$. Hence, in the finite grid world our agent will take finite time to either reach the goal or discover that it is impossible to reach the target. In worst case scenario our agent will traverse the entire maze. So for an $n * n$ maze, the maximum number of nodes it can explore is n^2 .

Part 2

In our Repeated Forward A* algorithm, when we resolve the ties between f-values using the g-values, it is observed that the selection of a node with a large g-value leads to the algorithm running comparatively faster when it comes to the execution time. For our algorithm we are selecting the node with minimum $f(n)$ value from the open list (priority queue) to explore and when the larger g-value is selected, than the heuristic value is in-turn smaller ($f(n) = g(n) + h(n)$). Choosing a larger g-value means we are choosing the node that is already closer to the goal since we always travel in the direction towards the goal unless the path in the direction to it is blocked, in which case we go around the blocked cell. Inversely, choosing a smaller g-value would take a long time from the agent's location to the goal since the agent would be checking multiple possible shortest paths. Therefore breaking ties in $f(n)$ with larger g-values leads to finding a path to the goal faster since we will be expanding less number of nodes.

```
***** A* Using Breaking Ties Between Fn in favour of Larger Gn values *****
START: 0 : 0
END: 100 : 100
Total Closed Nodes : 2967
AVG Runtime: 120253000 nanoseconds
***** A* Using Breaking Ties Between Fn in favour of Smaller Gn values *****
START: 0 : 0
END: 100 : 100
Total Closed Nodes : 108966
AVG Runtime: 5409975500 nanoseconds
```

Figure 2: A* run-time with larger $g(n)$ value against smaller $g(n)$ value

For 50 101*101 grids, implementing heap that breaks f-value ties in favor of larger g-values expanded on an average 3581 nodes but the heap in favor of smaller g-values expanded on an average 107512 nodes which is more by a factor of 30. And the run-time is more by a factor of about 60.

Part 3

After implementing and comparing both the algorithms we observed that Repeated Forward A* is faster compared to our Repeated Backward A* since the number of nodes expanded by the Backward A* is much larger than the number of nodes expanded by Forward A*. This is because the agent is closer to the Start, the Repeated Backward A* has less information, which necessitates fewer restarts

for the Repeated Forward A* while the Backward A* does not become aware of the wall/blockers until it comes near to the beginning.

For 50 mazes of 101x101, we can observe that the Forward A* Expands 5 times lesser nodes and runs 11 times after than Backward A*

```

Forward A* - Greater GN - AVG Expanded Nodes for 50 mazes : 3555
Forward A* - Greater GN - AVG Runtime for 50 mazes : 5.0325138E7
Backward A* - AVG Expanded Nodes for 50 mazes : 18879
Backward A* - AVG Runtime for 50 mazes : 5.70893124E8

```

Figure 3: Run-time for Forward A* and Backward A*

Part 4

Manhattan distance between any two cells in the grid is given by $|x1 - x2| + |y1 - y2|$ which is strictly based on vertical or horizontal paths. And similarly our agent can only move in 4 compass directions north (vertically up), south (vertically down), east (horizontally right), west (horizontally left) therefore Manhattan gives the faster path for our agent to reach the goal. Also since our agent cannot move diagonally we can say that the shortest path will always be found and Manhattan distance will always be consistent.

A heuristic function is called consistent, if the following holds true: $\forall(n, a, n') : h(n) \leq c(n, a, n') + h(n')$, where $c(n, a, n')$ is step cost for going from n to n' using action a . The project argues that “the Manhattan distances are consistent in grid-worlds in which the agent can move only in the four main compass directions.” In order to prove this, we will consider two cases. The first case is, the more obvious one, if the cell n is closer to the target cell than the cell n' . In this case, the $h(n) \leq h(n') \implies h(n) \leq h(n') + c(n, a, n')$ since it is located closer to the target cell. The second case is if the cell n' is located somewhere between the cell n and the target cell. In this case, $h(n') \leq h(n)$ since it is closer to target cell. However, since $c(n, a, n')$ is the cost for going from n to $n' \implies h(n) \leq h(n') + c(n, a, n')$ because after getting from n to n' , the rest of the path to the target cell cannot be more than $h(n')$. Therefore, the Manhattan distances are consistent.

Furthermore, because every time the Adaptive A* algorithm is ran and $h_{new}(s)$ is calculated, it updates the $h(s)$ not only of the cell where the agent is located, but for all of the cells that it expanded on its way to the target cell from the agent's cell. Therefore, the h-values $h_{new}(s)$ are not only admissible but also consistent.

Part 5

When we run both the algorithms Repeated Forward A* and Adaptive A* on 50 grids of 101 X 101 size, we can observe that the run-time of Adaptive A* is less, i.e., Adaptive A* finds the path faster when it is compared to Repeated Forward A*.

The reason behind it is that for Adaptive A*, the algorithm uses a modified heuristic which is calculated as a difference between the goal's $g(n)$ and the current node's $g(n)$. This leads to a more realistic heuristic value that is admissible and closer to the actual heuristic.

The Adaptive A* algorithm does not expand as many cells as the Repeated Forward A* because each time it finds the path to the target cell, it updates the $h(n)$ of cells in the closed list. Therefore, when the agent follows that path, runs into a blocked cell, and reruns the algorithm to find a new path, then the algorithm would use the updated $h(n)$ from the previous iteration of A*.

In the above run, we can observe that Adaptive A* expands 5% lesser nodes than Forward A*

```
Forward A* - Greater GN - AVG Expanded Nodes for 50 mazes : 3515
Forward A* - Greater GN - AVG Runtime for 50 mazes : 7.098361E7
Adaptive A* - AVG Expanded Nodes for 50 mazes : 3345
Adaptive A* - AVG Runtime for 50 mazes : 7.073805E7
```

Figure 4: Run-time for Forward A* and Adaptive A*

Part 6

In our implementation we are using 2-D array-list for maze and have created a node class for the agent to follow to reach the goal. The node class includes attributes like coordinates of the node stored as integer array, the parent of the node as a pointer and f, g, h values for the respective node. This class in total uses 40 Bytes of memory.

Now in order to operate within 4 MBytes, we would have to be using a maze of size not more than 308×308 . In order to calculate the size, we used the following equation: $2 * n^2 + (n^2 * 40) = 4 * 10^6$, where n is the number of cells, 40 is the number of Bytes needed per each cell in our implementation, and $4 * 10^6$ is the 4 MBytes represented in Bytes. The first $2 * n^2$ represents the size of 2-D array of the maze, then $n^2 * 40$ corresponds to the number of nodes time the size in Bytes of the node structure. As a result of solving this equation for n , we get that $n = \sqrt{(4 * 10^6 / 42)} \implies n = 308.6$. Therefore, the largest maze size that we could operate on within the memory limit of 4 MBytes is 308×308 .